

# Assignment 4: Predicting Automobile Pricing Using Neural Networks

Nabeel Khan, Nha Nguyen, Razzaq

2022-12-02

## R Markdown

This is an R Markdown document. Markdown is a simple formatting syntax for authoring HTML, PDF, and MS Word documents. For more details on using R Markdown see <http://rmarkdown.rstudio.com>.

When you click the **Knit** button a document will be generated that includes both content as well as the output of any embedded R code chunks within the document. You can embed an R code chunk like this:

```
library(readxl)
library(ISLR)
library(caret)
```

```
## Loading required package: ggplot2
```

```
## Loading required package: lattice
```

```
library(pROC)
```

```
## Type 'citation("pROC")' for a citation.
```

```
##
```

```
## Attaching package: 'pROC'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      cov, smooth, var
```

```
library(ROSE)
```

```
## Loaded ROSE 0.0-4
```

```
library(robustbase)
library(smotefamily)
library(ROCR)
library(rpart)
library(rpart.plot)
library(ggplot2)
library(neuralnet)
```

```
##
## Attaching package: 'neuralnet'

## The following object is masked from 'package:ROCR':
##
##      prediction
```

```
library(GGally)
```

```
## Registered S3 method overwritten by 'GGally':
##   method from
##   +.gg      ggplot2
```

```
library(corrr)
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1.3.2 --
```

```
## v tibble  3.1.8      v dplyr   1.0.10
## v tidyr   1.2.1      v stringr 1.4.1
## v readr   2.1.2      v forcats 0.5.2
## v purrr   0.3.4
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::compute() masks neuralnet::compute()
## x dplyr::filter()  masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## x purrr::lift()    masks caret::lift()
```

```
library(ROSE)
```

```
#importing the dataset
Assi4 <- read_excel("Assi4.xlsx")
#Summary of the data present in Assi4
summary(Assi4)
```

```
##      Price      Age      KM      HP
##  Min.   :13000  Min.   :23.00  Min.   : 9752  Min.   :100.0
## 1st Qu.:15800  1st Qu.:26.00  1st Qu.:25448  1st Qu.:100.0
## Median :16300  Median :28.00  Median :32222  Median :113.0
## Mean   :17251  Mean   :27.94  Mean   :33628  Mean   :126.9
## 3rd Qu.:18000  3rd Qu.:30.00  3rd Qu.:42210  3rd Qu.:113.0
## Max.   :22800  Max.   :33.00  Max.   :67662  Max.   :195.0
##      MC      Colour      Auto      CC
##  Min.   :0.0000  Length:31      Min.   :0.00000  Min.   :1400
## 1st Qu.:0.0000  Class :character 1st Qu.:0.00000  1st Qu.:1400
## Median :1.0000  Mode  :character Median :0.00000  Median :1600
## Mean   :0.6452                      Mean   :0.03226  Mean   :1574
## 3rd Qu.:1.0000                      3rd Qu.:0.00000  3rd Qu.:1600
## Max.   :1.0000                      Max.   :1.00000  Max.   :1800
##      Grs      Wght      G_P      Mfr_G
##  Min.   :5.000  Min.   :1069  Min.   : 4.000  Min.   :0.0000
```

```
## 1st Qu.:5.000 1st Qu.:1104 1st Qu.: 4.000 1st Qu.:1.0000
## Median :5.000 Median :1124 Median : 4.000 Median :1.0000
## Mean :5.194 Mean :1130 Mean : 4.903 Mean :0.9032
## 3rd Qu.:5.000 3rd Qu.:1149 3rd Qu.: 4.000 3rd Qu.:1.0000
## Max. :6.000 Max. :1189 Max. :20.000 Max. :1.0000
## Abag_2 AC Comp CD
## Min. :0.0000 Min. :0.0000 Min. :0.0000 Min. :0.0000
## 1st Qu.:1.0000 1st Qu.:0.0000 1st Qu.:1.0000 1st Qu.:0.0000
## Median :1.0000 Median :1.0000 Median :1.0000 Median :1.0000
## Mean :0.9355 Mean :0.5484 Mean :0.9032 Mean :0.6129
## 3rd Qu.:1.0000 3rd Qu.:1.0000 3rd Qu.:1.0000 3rd Qu.:1.0000
## Max. :1.0000 Max. :1.0000 Max. :1.0000 Max. :1.0000
## Clock Pw Radio SpM
## Min. :0.0000 Min. :0.0000 Min. :0.00000 Min. :0.0000
## 1st Qu.:1.0000 1st Qu.:1.0000 1st Qu.:0.00000 1st Qu.:1.0000
## Median :1.0000 Median :1.0000 Median :0.00000 Median :1.0000
## Mean :0.9032 Mean :0.9032 Mean :0.06452 Mean :0.8065
## 3rd Qu.:1.0000 3rd Qu.:1.0000 3rd Qu.:0.00000 3rd Qu.:1.0000
## Max. :1.0000 Max. :1.0000 Max. :1.00000 Max. :1.0000
## M_Rim Tow_Bar
## Min. :0.0000 Min. :0.0000
## 1st Qu.:0.0000 1st Qu.:0.0000
## Median :0.0000 Median :0.0000
## Mean :0.3871 Mean :0.1613
## 3rd Qu.:1.0000 3rd Qu.:0.0000
## Max. :1.0000 Max. :1.0000
```

```
# Check missing values
sapply(Assi4,function(x) sum(is.na(x))) #There are no missing values in all columns.
```

```
## Price Age KM HP MC Colour Auto CC Grs Wght
## 0 0 0 0 0 0 0 0 0 0
## G_P Mfr_G Abag_2 AC Comp CD Clock Pw Radio SpM
## 0 0 0 0 0 0 0 0 0 0
## M_Rim Tow_Bar
## 0 0
```

```
# Factor categorical variables
Assi4$Colour <- as.factor(Assi4$Colour)
#Attach Assi4 to data1
data1 <- Assi4
# Calculate correlation coefficient. #Factors that influence a customer's decision
#to buy a car are Horsepower (HP), Cylinder volume in cubic cms (CC) and Weight (Wght)
#whose correlation with car price are 0.923, 0.890 and 0.856 accordingly.

correlate(data1)
```

```
## Non-numeric variables removed from input: 'Colour'
## Correlation computed with
## * Method: 'pearson'
## * Missing treated using: 'pairwise.complete.obs'
```

```
## # A tibble: 21 x 22
```

```
##      term      Price      Age      KM      HP      MC      Auto      CC      Grs
##      <chr>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>
## 1 Price NA          0.273 -0.0733  0.923 -0.117 -0.0183  0.890  0.760
## 2 Age  0.273 NA          0.0356  0.164  0.347  0.127  0.285  0.0381
## 3 KM   -0.0733  0.0356 NA          -0.0601 -0.0701  0.137  0.0738 -0.111
## 4 HP   0.923  0.164 -0.0601 NA          -0.223 -0.0682  0.892  0.896
## 5 MC   -0.117  0.347 -0.0701 -0.223 NA          -0.246 -0.127 -0.319
## 6 Auto -0.0183  0.127  0.137 -0.0682 -0.246 NA          0.0314 -0.0894
## 7 CC   0.890  0.285  0.0738  0.892 -0.127  0.0314 NA          0.737
## 8 Grs  0.760  0.0381 -0.111  0.896 -0.319 -0.0894  0.737 NA
## 9 Wght 0.856  0.248 -0.0122  0.918 -0.208  0.226  0.846  0.815
## 10 G_P 0.0369 -0.160 -0.261  0.00685  0.0849 -0.0528  0.0497  0.0152
## # ... with 11 more rows, and 13 more variables: Wght <dbl>, G_P <dbl>,
## #   Mfr_G <dbl>, Abag_2 <dbl>, AC <dbl>, Comp <dbl>, CD <dbl>, Clock <dbl>,
## #   Pw <dbl>, Radio <dbl>, SpM <dbl>, M_Rim <dbl>, Tow_Bar <dbl>
```

```
# Test relationship between categorical and numerical variables using Chi-squared test
chisq.test(data1$Price, data1$Colour)
```

```
## Warning in chisq.test(data1$Price, data1$Colour): Chi-squared approximation may
## be incorrect
```

```
##
## Pearson's Chi-squared test
##
## data: data1$Price and data1$Colour
## X-squared = 99.338, df = 90, p-value = 0.235
```

- (1) After your EDA, what factors do you think influence a customer's decision to buy a car? What are the objectives of the model that Farid plans to build?

#### Question 1 Answer

There are some variables present in the excel sheet where the variables take the same value for all the observations. These variables can directly be removed from the excel sheet as they do not make any difference to our models as they have the same value for all observations. These variables are - Fuel: (Fuel type (Petrol, Diesel, CNG)), Drs(Number of doors), Cyl(Number of cylinders), ABS(Anti-lock brake system (Yes=1, No=0)), Abag1(Driver airbag (Yes=1, No=0)), PStr(Power steering (Yes=1, No=0)). Remaining variables are the ones that influence a customer's decision to buy a car but it can be noted that some variables are more correlated to variable 'Price' than other variables. We have shown this in our code. Although some variables are more correlated than other to price, we have decided to use all the variables except for the ones removed in the first paragraph

We have then performed normalization to the data after importing it bringing it to a scale between 0 and 1. This makes it easier for neural network model.

Objectives of the model that Farid plans to build: For predicting the prices, Farid decided that he would try various machine-learning methods, including linear regression. He also knew that neural networks excelled in predicting price, so he decided to use a feed-forward neural network to train data and accurately predict the price. He wondered how neural networks would compare to linear regression. For long-term marketing, he wanted to decide on one particular computing system to determine prices.

```
# (2) Construct a neural network model. Validate and interpret the model
# using a different number of hidden neurons.
```

```
#Create a function to normalize data between 0 and 1
```

```
mynormalization <- function(x)
{
  (x - min(x))/ (max(x)-min(x))
}
```

```
#Apply normalization to only those columns that are numerical
```

```
data1 <- Assi4%>% mutate_if(is.numeric, mynormalization)
summary(data1)
```

```
##      Price      Age      KM      HP
##  Min.   :0.0000  Min.   :0.0000  Min.   :0.0000  Min.   :0.0000
## 1st Qu.:0.2857  1st Qu.:0.3000  1st Qu.:0.2710  1st Qu.:0.0000
## Median :0.3367  Median :0.5000  Median :0.3880  Median :0.1368
## Mean   :0.4338  Mean   :0.4935  Mean   :0.4123  Mean   :0.2832
## 3rd Qu.:0.5102  3rd Qu.:0.7000  3rd Qu.:0.5605  3rd Qu.:0.1368
## Max.   :1.0000  Max.   :1.0000  Max.   :1.0000  Max.   :1.0000
##      MC      Colour      Auto      CC
##  Min.   :0.0000  Black : 5  Min.   :0.00000  Min.   :0.0000
## 1st Qu.:0.0000  Blue  : 9  1st Qu.:0.00000  1st Qu.:0.0000
## Median :1.0000  Green : 1  Median :0.00000  Median :0.5000
## Mean   :0.6452  Grey  :12  Mean   :0.03226  Mean   :0.4355
## 3rd Qu.:1.0000  Red   : 2  3rd Qu.:0.00000  3rd Qu.:0.5000
## Max.   :1.0000  Silver: 2  Max.   :1.00000  Max.   :1.0000
##      Grs      Wght      G_P      Mfr_G
##  Min.   :0.0000  Min.   :0.0000  Min.   :0.00000  Min.   :0.0000
## 1st Qu.:0.0000  1st Qu.:0.2917  1st Qu.:0.00000  1st Qu.:1.0000
## Median :0.0000  Median :0.4583  Median :0.00000  Median :1.0000
## Mean   :0.1935  Mean   :0.5081  Mean   :0.05645  Mean   :0.9032
## 3rd Qu.:0.0000  3rd Qu.:0.6667  3rd Qu.:0.00000  3rd Qu.:1.0000
## Max.   :1.0000  Max.   :1.0000  Max.   :1.00000  Max.   :1.0000
##      Abag_2      AC      Comp      CD
##  Min.   :0.0000  Min.   :0.0000  Min.   :0.0000  Min.   :0.0000
## 1st Qu.:1.0000  1st Qu.:0.0000  1st Qu.:1.0000  1st Qu.:0.0000
## Median :1.0000  Median :1.0000  Median :1.0000  Median :1.0000
## Mean   :0.9355  Mean   :0.5484  Mean   :0.9032  Mean   :0.6129
## 3rd Qu.:1.0000  3rd Qu.:1.0000  3rd Qu.:1.0000  3rd Qu.:1.0000
## Max.   :1.0000  Max.   :1.0000  Max.   :1.0000  Max.   :1.0000
##      Clock      Pw      Radio      SpM
##  Min.   :0.0000  Min.   :0.0000  Min.   :0.00000  Min.   :0.0000
## 1st Qu.:1.0000  1st Qu.:1.0000  1st Qu.:0.00000  1st Qu.:1.0000
## Median :1.0000  Median :1.0000  Median :0.00000  Median :1.0000
## Mean   :0.9032  Mean   :0.9032  Mean   :0.06452  Mean   :0.8065
## 3rd Qu.:1.0000  3rd Qu.:1.0000  3rd Qu.:0.00000  3rd Qu.:1.0000
## Max.   :1.0000  Max.   :1.0000  Max.   :1.00000  Max.   :1.0000
##      M_Rim      Tow_Bar
##  Min.   :0.0000  Min.   :0.0000
## 1st Qu.:0.0000  1st Qu.:0.0000
## Median :0.0000  Median :0.0000
## Mean   :0.3871  Mean   :0.1613
## 3rd Qu.:1.0000  3rd Qu.:0.0000
```

```
## Max. :1.0000 Max. :1.0000
```

```
#Create training and test data
```

```
set.seed(124)
```

```
indx <- sample(2, nrow(data1), replace = T, prob = c(0.7, 0.3))
```

```
train <- data1[indx == 1, ]
```

```
test <- data1[indx == 2, ]
```

```
#Create Neural Network Model
```

```
set.seed(12321)
```

```
nnModel1 <- neuralnet(Price ~ Age + KM + HP + MC + Auto + CC + Grs + Wght + G_P + Mfr_G + Abag_2 +  
AC + Comp + CD + Clock + Pw + Radio + SpM + M_Rim + Tow_Bar, data = train,  
err.fct = 'sse', hidden = c(3,2), threshold = 0.01, linear.output = T)
```


```
#Check for error and other parameters using result.matrix
```

```
nnModel1$result.matrix
```

```
##                                     [,1]  
## error                             0.011562574  
## reached.threshold                 0.004938142  
## steps                             111.000000000  
## Intercept.to.1layhid1             0.820224401  
## Age.to.1layhid1                   3.768349990  
## KM.to.1layhid1                    -1.789664164  
## HP.to.1layhid1                   0.386294119  
## MC.to.1layhid1                   0.576393249  
## Auto.to.1layhid1                 0.467906031  
## CC.to.1layhid1                   1.066735548  
## Grs.to.1layhid1                  -0.851800712  
## Wght.to.1layhid1                 0.803473099  
## G_P.to.1layhid1                  -1.882545614  
## Mfr_G.to.1layhid1                -0.864612567  
## Abag_2.to.1layhid1               -1.644053347  
## AC.to.1layhid1                   -1.569773474  
## Comp.to.1layhid1                 -1.007856615  
## CD.to.1layhid1                   0.231672799  
## Clock.to.1layhid1                1.042445234  
## Pw.to.1layhid1                   -1.585747574  
## Radio.to.1layhid1                1.002211915  
## SpM.to.1layhid1                  -0.388048569  
## M_Rim.to.1layhid1                1.705002655  
## Tow_Bar.to.1layhid1              -0.371200763  
## Intercept.to.1layhid2             0.708950973  
## Age.to.1layhid2                  -0.212681901  
## KM.to.1layhid2                   -3.668135300  
## HP.to.1layhid2                   -3.979322733  
## MC.to.1layhid2                   -0.450550977  
## Auto.to.1layhid2                 -0.666357889  
## CC.to.1layhid2                   -3.296567317  
## Grs.to.1layhid2                  -5.158770432  
## Wght.to.1layhid2                 -1.289448615  
## G_P.to.1layhid2                  0.102845964  
## Mfr_G.to.1layhid2                1.674982481
```

```
## Abag_2.to.1layhid2      -1.321643395
## AC.to.1layhid2         -4.150138038
## Comp.to.1layhid2       -0.379616557
## CD.to.1layhid2         -6.717130530
## Clock.to.1layhid2      -0.729597629
## Pw.to.1layhid2         -1.037170271
## Radio.to.1layhid2      3.968872090
## SpM.to.1layhid2        -0.453253314
## M_Rim.to.1layhid2      -6.412707410
## Tow_Bar.to.1layhid2    -2.491349543
## Intercept.to.1layhid3  -0.655843480
## Age.to.1layhid3        1.068465592
## KM.to.1layhid3         -5.050540248
## HP.to.1layhid3         7.462586006
## MC.to.1layhid3         -2.434637303
## Auto.to.1layhid3       0.387411557
## CC.to.1layhid3         2.489236595
## Grs.to.1layhid3        5.192727882
## Wght.to.1layhid3       1.825247555
## G_P.to.1layhid3        -7.026036540
## Mfr_G.to.1layhid3      -1.515640751
## Abag_2.to.1layhid3     0.693283328
## AC.to.1layhid3         1.256206324
## Comp.to.1layhid3       -0.242563934
## CD.to.1layhid3         0.353113843
## Clock.to.1layhid3      -0.465777167
## Pw.to.1layhid3         -1.470714639
## Radio.to.1layhid3      -5.119614668
## SpM.to.1layhid3        -1.488639744
## M_Rim.to.1layhid3      -0.660696697
## Tow_Bar.to.1layhid3    -2.204517686
## Intercept.to.2layhid1  0.306985563
## 1layhid1.to.2layhid1   0.364099760
## 1layhid2.to.2layhid1   -1.531865567
## 1layhid3.to.2layhid1   3.146702685
## Intercept.to.2layhid2  -0.217887538
## 1layhid1.to.2layhid2   -2.530063358
## 1layhid2.to.2layhid2   1.904271813
## 1layhid3.to.2layhid2   -0.095498817
## Intercept.to.Price     -0.100395726
## 2layhid1.to.Price      1.134773107
## 2layhid2.to.Price      -0.762774759
```

```
#Error on training set for this model is 0.011562574
#Plot the model
plot(nnModel1, rep = 'best')
```



```
# Predict on test data
pr.nn <- neuralnet::compute(nnModel1, test[,1:22])

# Compute mean squared error
pr.nn_ <- pr.nn$net.result * (max(data1$Price) - min(data1$Price)) + min(data1$Price)
pr.nn_
```

```
##           [,1]
## [1,] 0.6309953
## [2,] 0.3488842
## [3,] 0.4015929
## [4,] 0.3336059
## [5,] 0.3020155
## [6,] 0.4583227
## [7,] 0.2756727
## [8,] 0.2611084
## [9,] 0.5216279
## [10,] 0.1725715
```

```
test.r <- (test$Price) * (max(data1$Price) - min(data1$Price)) +
  min(data1$Price)
test.r
```

```
## [1] 0.3877551 0.3061224 0.4081633 0.3061224 0.4081633 0.2857143 0.4081633
## [8] 0.3061224 0.2857143 0.1020408
```

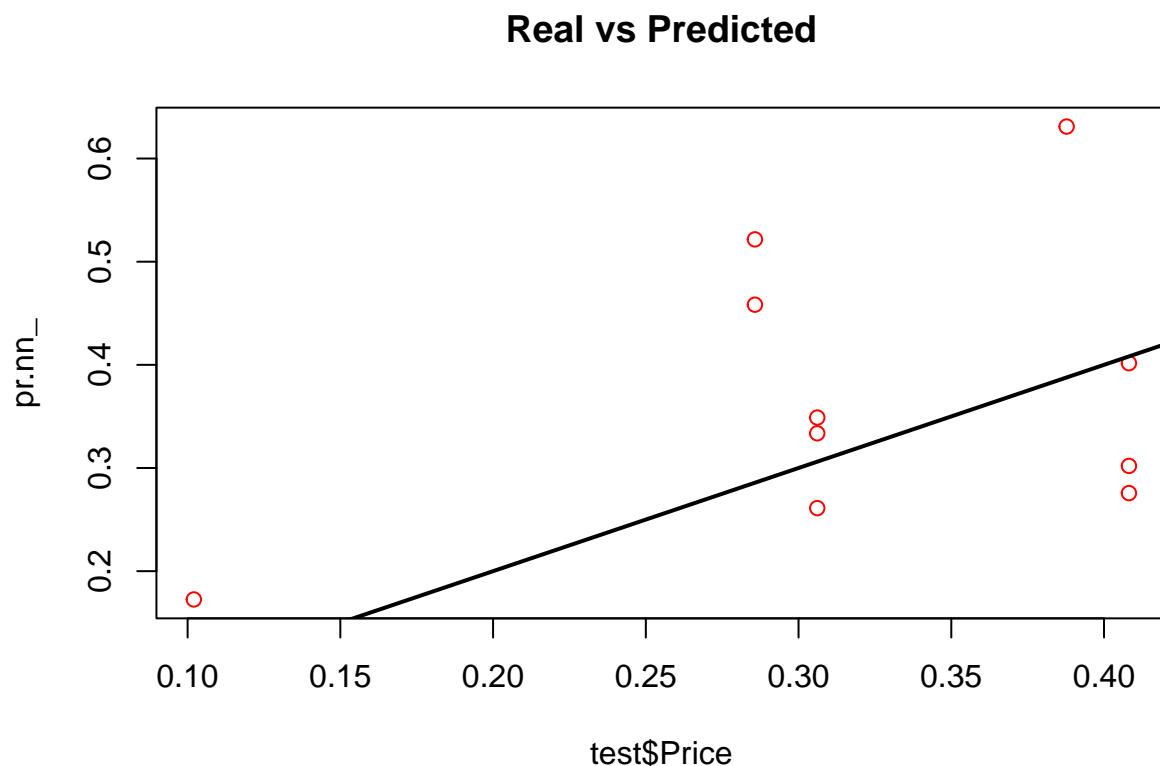


```
#MSE.nn is the Mean Squared Error on the test data
MSE.nn <- sum((test.r - pr.nn_)^2) / nrow(test)
MSE.nn
```

```
## [1] 0.01830637
```

```
#MSE.nn is the mean squared error for test data which is 0.01830637
```

```
# Plot regression line
plot(test$Price, pr.nn_, col = "red",
      main = 'Real vs Predicted')
abline(0, 1, lwd = 2)
```



I have used various hidden neurons like (5,4,2), (4,2), (2,1) and (3,2) and have finally decided to use (3,2) after comparing the error for each on training and test data. For the Real vs pred graph, we can see that the predictions (red circles) made by the neural network are in general concentrated around the line (a perfect alignment with the line would indicate an MSE of 0 and thus an ideal prediction).

```
# (3) Compare your neural network models with linear regression model. Which one is better?
```

```
# Converting "Price" as numeric and factorize "Colour" variable
data1$Price <- as.numeric(data1$Price)
data1$Colour <- as.factor(data1$Colour)
```

```
# Split train test and test set for lm regression
set.seed(123)
```

```
lm.indx <- sample(2,nrow(data1), replace = T, prob = c(0.7,0.3))
lm.train <- data1[lm.indx == 1,]
lm.test <- data1[lm.indx ==2,]

full <- lm(lm.train$Price ~ .,data = lm.train)
null <- lm(lm.train$Price ~ 1,data = lm.train)

step(null, scope = list(lower = null, upper = full), direction = "forward", trace = 0)
```

```
##
## Call:
## lm(formula = lm.train$Price ~ HP + Colour + SpM + Mfr_G + Clock,
##     data = lm.train)
##
## Coefficients:
## (Intercept)          HP    ColourBlue    ColourGreen    ColourGrey
##    0.209212    0.604399   -0.005046     0.021214     0.043884
##   ColourRed  ColourSilver          SpM          Mfr_G          Clock
##  -0.190328   -0.145074   -0.092501     0.072009     0.075781
```

#### *#Create LM Model*

*#I have taken -Tow\_Bar - Radio -Pw -Comp -Abag\_2 -G\_P -Colour variables out as when we split it into train and test data, that is because the datasets were imbalanced and linear regression failed. You might think this is wrong to remove variables like this but I can justify it as they did not have a high correlation with 'price' when we ended up doing EDA in question 1. Hence it is justified.*

```
lmModel <- lm(formula = lm.train$Price ~ . -Tow_Bar - Radio -Pw -Comp -Abag_2
              -G_P -Colour, data = lm.train)
summary(lmModel)
```

```
##
## Call:
## lm(formula = lm.train$Price ~ . - Tow_Bar - Radio - Pw - Comp -
##     Abag_2 - G_P - Colour, data = lm.train)
##
## Residuals:
##      1      2      3      4      5      6      7
## -1.681e-02 -8.041e-02  9.722e-02  2.373e-17  2.503e-17 -4.913e-17 -5.329e-18
##      8      9     10     11     12     13     14
## -6.306e-03  7.687e-02 -2.186e-02  1.681e-02 -7.564e-02  1.013e-02  6.514e-02
##     15     16     17     18     19     20
##  4.019e-02 -5.347e-17 -1.013e-02 -6.712e-02 -2.808e-02 -4.783e-17
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.24447    0.13832   1.767   0.137
## Age          0.11052    0.12283   0.900   0.409
## KM          -0.14686    0.21309  -0.689   0.521
## HP           0.14364    0.58200   0.247   0.815
## MC          -0.07898    0.08579  -0.921   0.400
## Auto        -0.16712    0.24398  -0.685   0.524
```

```
## CC          0.35847    0.48523    0.739    0.493
## Grs         -0.19703    0.15351   -1.284    0.256
## Wght        0.39744    0.45526    0.873    0.423
## Mfr_G       -0.09667    0.14832   -0.652    0.543
## AC          -0.18595    0.19272   -0.965    0.379
## CD          -0.01800    0.06482   -0.278    0.792
## Clock       0.16634    0.12167    1.367    0.230
## SpM        -0.09030    0.08445   -1.069    0.334
## M_Rim       0.04233    0.09388    0.451    0.671
##
## Residual standard error: 0.08943 on 5 degrees of freedom
## Multiple R-squared:  0.9616, Adjusted R-squared:  0.8541
## F-statistic: 8.948 on 14 and 5 DF,  p-value: 0.01219
```

```
# LR prediction on train data and test data
```

```
train_MSE = mean(lmModel$residuals^2)
test_MSE = mean((lm.test$Price - predict.lm(lmModel, lm.test)) ^ 2)
train_MSE #0.001999398
```

```
## [1] 0.001999398
```

```
test_MSE #0.01692538
```

```
## [1] 0.01692538
```

```
fitted(lmModel)
```

```
##          1          2          3          4          5          6          7          8
## 0.8331322 0.7589858 0.8262493 1.0000000 0.3877551 0.4081633 0.4081633 0.3124286
##          9         10         11         12         13         14         15         16
## 0.3312929 0.3585925 0.4469698 0.3613524 0.3980372 0.2205757 0.2659326 0.2602041
##         17         18         19         20
## 0.3162485 0.2712032 0.3137951 0.1836735
```

```
coefficients(lmModel)
```


```
## (Intercept)      Age      KM      HP      MC      Auto
## 0.24447486 0.11051556 -0.14685827 0.14363655 -0.07897997 -0.16712447
##          CC      Grs      Wght      Mfr_G      AC      CD
## 0.35847007 -0.19702882 0.39743509 -0.09666580 -0.18595028 -0.01799818
##      Clock      SpM      M_Rim
## 0.16633640 -0.09029839 0.04232620
```

```
residuals((lmModel))
```


```
##          1          2          3          4          5
## -1.680568e-02 -8.041438e-02 9.722006e-02 2.372802e-17 2.502907e-17
##          6          7          8          9         10
## -4.913036e-17 -5.328593e-18 -6.306164e-03 7.687033e-02 -2.185778e-02
##         11         12         13         14         15
## 1.680568e-02 -7.563812e-02 1.012606e-02 6.513861e-02 4.018986e-02
##         16         17         18         19         20
## -5.346717e-17 -1.012606e-02 -6.712160e-02 -2.808081e-02 -4.782932e-17
```

```
plot(lmModel)
```


```
## Warning: not plotting observations with leverage one:  
## 4, 5, 6, 7, 16, 20
```




lm(lm.train\$Price ~ . - Tow\_Bar - Radio - Pw - Comp - Abag\_2 - G\_P - Colour ...)



lm(lm.train\$Price ~ . - Tow\_Bar - Radio - Pw - Comp - Abag\_2 - G\_P - Colour ...



Fitted values  
 $\text{lm}(\text{lm.train}\$Price \sim . - \text{Tow\_Bar} - \text{Radio} - \text{Pw} - \text{Comp} - \text{Abag\_2} - \text{G\_P} - \text{Colour} \dots)$



`lm(lm.train$Price ~ . - Tow_Bar - Radio - Pw - Comp - Abag_2 - G_P - Colour ...`

As we can see, I have got error MSE on train and test data of linear regression model in this question and also neural network model in the previous question. It is clearly seen that error is less for linear regression model when compared to neural network model but it is noteworthy to keep in mind that we considered all the variables for our neural network model. Also, the training and test error for neural network model were almost similar whereas for LM model we can clearly see that training set has very less error when compared to test error. This can mean that our LM model can be slightly overfitted. As to which one is better, I would say it depends on the individual deciding as we can prune both models to satisfy our needs with the question but I personally would go for the neural network model as it would lead to more accurate and faster results when we are working with huge and changing datasets.

Question 4 (4) Make a decision and offer your recommendations. Decision: Neural network is a better model  
Recommendations: The dataset used for this exercise is very small when compared to actual practical datasets for car prices. For practical use, neural network is the better model as it takes into account all the variables and can easily change accordingly if our needs from the model change in the future. I would still like to point out that as we are getting lower error for LM model, it would not be fair to completely neglect that option.

““